ANON POLITICIAN

Master Vision, Research Specification and Handoff to the Next Model

Version 1.0 | 14 October 2026

A proposed AI-assisted democratic policy intelligence system for Aotearoa New Zealand. AI supports reasoning; elected humans remain accountable.

PRIVATE RESEARCH HANDOFF
Includes the original source transcript and prior technical artifacts. Review identifying context before sharing publicly.

Navigation

- Start here: what the next model needs to understand
- 1. The founder's problem statement and intent
- 2. Product vision
- 3. The core democratic loop
- 4. Recurring citizen polls and consultation
- 5. The opposing-party adversarial model
- 6. Forecasting and future trajectories
- 7. CASCADE: knowledge and institutional memory
- 8. AURA: values, methods, and original metrics
- 9. Original DEEP idea and ecological accounting
- 10. New Zealand context and sources
- 11. Design principles and safety invariants
- 12. Main risks to investigate
- 13. First prototype
- 14. Evaluation plan
- 15. Current implementation state
- 16. Staged development plan
- 17. Open design decisions
- 18. Working instructions for the next model
- 19. Existing artifacts and privacy warning

Important epistemic rule: coherence is not truth; test results must not be reported beyond what they actually establish.

Start here: what the next model needs to understand

Anon Politician is intended to become an **AI-assisted democratic policy intelligence system**. The central idea is not for AI to govern, decide policy, or tell citizens how to vote. It is to help elected representatives and the public develop, challenge, compare, improve, and track policies using public input, evidence, long-term scenario analysis, structured institutional memory, and serious adversarial scrutiny.

The founder's concern is that many people in New Zealand feel their ability to forge a secure, meaningful future has weakened, while political competition can become so adversarial that parties resist ideas because they originated with the opposing side. The project should investigate that concern rather than assume every part of it is already proven.

The founder's distinctive proposed mechanism is a **recurring citizen-input loop**. Politicians could publish policy questions on which they want public feedback; New Zealanders could participate in regular, potentially monthly polls or consultations; and AI could analyse the feedback alongside evidence, alternative policies, forecasts, and critiques from opposing perspectives. Representatives would then be able to explain what they changed, what they did not change, and why. Later, the system could compare expectations with real outcomes.

The desired incentive is to make responsive, transparent, evidence-informed leadership competitively valuable: leaders can demonstrate that they listen, examine trade-offs, improve proposals, and deliver results. This is a hypothesis to test, not a guaranteed effect. Monthly polls are not automatically representative or binding, and public preference is not identical to evidence, rights, or long-term public interest.

The technical origin is the user's AURA Protocol and Pyramid Cascade Architecture, also referred to as CASCADE. The founding idea is to keep knowledge and policy reasoning traceable, preserve contradictions, and review conclusions when supporting foundations change—without silently erasing ideas or claiming that coherence proves truth.

Immediate next step: choose one tractable New Zealand policy question; build a small end-to-end prototype; compare it against ordinary policy analysis and a simpler evidence ledger; and evaluate whether it measurably improves traceability, understanding, forecast calibration, balanced critique, or policy-option quality.

1. The founder's problem statement and intent

The founder's motivating view, expressed in conversation, is that New Zealand has stripped or failed to preserve enough people's ability to forge a future here. The concern is broad and can include the practical ability to build secure lives, access opportunity, plan for the long term, and feel that political decisions respond to the public.

A second concern is the structure of partisan competition. Labour, National, and other political groupings have distinct values and policy theories; that diversity is legitimate. But if parties treat an idea as unacceptable merely because it came from an opponent, useful proposals may be dismissed before their evidence or design is tested. The project should help distinguish **disagreement about values**, **disagreement about facts**, **different forecasts**, **different interests**, and **identity-based rejection**.

The goal is not forced consensus or ideological convergence. The goal is constructive contestability: proposals face serious opposition, the proposer can respond, evidence and trade-offs are made visible, and decision-makers remain accountable for their choices.

The project must investigate the founder's assumptions as carefully as it investigates those of political institutions. For example, it must remain open to the possibility that public opinion is mixed or uninformed on a particular technical question, that parties reject some proposals for substantive reasons, that some

disagreements are legitimate value conflicts, and that existing government processes already do some parts of this work well.

2. Product vision

2.1 Product in one sentence

An independent AI-assisted policy workspace that connects citizen input, evidence, cross-perspective critique, scenario forecasts, policy revision, and outcome tracking—while leaving political decisions and legal accountability with humans.

2.2 What it should help people do

- Let elected representatives ask clear questions about policy options and invite public input on a recurring basis.
- Let citizens state priorities, compare alternatives, explain concerns, and see how their views were considered.
- Build an inspectable record of each policy proposal: its goals, assumptions, evidence, dependencies, uncertainties, affected groups, alternatives, and decision history.
- Use AURA-inspired value and action principles without pretending moral disagreements can all be reduced to one objective score.
- Use CASCADE to identify which claims and conclusions need review when sources, assumptions, or dependencies change.
- Use adversarial models to produce robust, fair critiques from different policy perspectives—including a proposed role for a model configured or trained with input from an opposing party.
- Use forecasting and scenario analysis to examine plausible present and future outcomes, with uncertainty and limitations stated explicitly.
- Generate improvements and alternatives instead of merely saying yes or no.
- Publish a decision trace showing which public input and evidence were considered, which trade-offs remained, what the representative decided, and why.
- Track what happened after implementation and compare outcomes with prior forecasts and stated objectives.

2.3 What it is not

- Not an AI government, automated lawmaker, or substitute for elected representatives.
- Not an electoral prediction engine, voting-advice service, candidate ranking, party leaderboard, or personalised political persuasion system.
- Not a system that equates a poll majority with truth, lawfulness, fairness, or the interests of every person affected.
- Not a mechanism for automatically deleting a policy or claim because it conflicts with a core value, party position, or model output.
- Not a promise that an AI model can discover one universally correct definition of the national interest.
- Not a claim that the AURA metrics or CASCADE architecture have already been independently validated.

3. The core democratic loop

The proposed workflow is a recurring loop. The first implementation should be small and transparent; each step should leave a record that can be reviewed later.

1. **A question is opened.** A representative, committee, civic group, or authorised public process publishes a neutral question about a real policy decision. The question includes the decision context, options under consideration, constraints, and what feedback can realistically influence.
2. **Citizens contribute.** New Zealanders can vote on options, rank priorities, express uncertainty, explain concerns, submit evidence, or propose alternatives. Participation should be accessible and privacy-

conscious. A poll must not be presented as national consensus if the sample is self-selected or otherwise unrepresentative.

3. **The system characterises the input.** It reports response counts, recruitment and sampling methods, missing groups, non-response, uncertainty, geographic and demographic skews where lawful and appropriate, and the distinction between a consultation result and a representative survey. It summarises themes without treating the AI summary as a substitute for the original submissions.
4. **Evidence is assembled.** Claims are linked to sources, dates, methods, source quality, limitations, conflicts of interest where known, and the extent to which evidence applies to New Zealand.
5. **CASCADE structures reasoning.** Each important factual claim, forecast, assumption, value commitment, and policy option is stored with provenance and dependencies. Foundational changes create a review queue rather than an automatic verdict.
6. **Adversarial review occurs.** One or more models challenge the proposal from different perspectives. The challenge should steelman the strongest opposing case, identify overlooked groups and implementation risks, and use verifiable sources rather than manufacture arguments.
7. **Forecasts and alternatives are compared.** Models explore plausible scenarios for the proposal, alternatives, and the status quo. Forecasts show assumptions, ranges, uncertainty, affected groups, fiscal and non-fiscal impacts, and what new evidence would change the result.
8. **A revision round occurs.** The proposing representative or policy team can respond to the critique, correct factual errors, revise the policy, or explain why an objection does not change their position. The system retains both the initial and revised versions.
9. **The public receives a decision trace.** The report shows what participants said, how the results should be interpreted, what evidence was used, which objections were raised, what changed, what did not change, and why. It records unresolved disagreement rather than claiming unanimity.
10. **The decision-maker remains accountable.** A human authority decides within existing democratic, legal, and constitutional processes. The record identifies the responsible decision-maker and the reasons given.
11. **Outcomes are reviewed.** After implementation, the system tracks agreed indicators, checks forecasts against observed outcomes, records unintended consequences, and reopens dependent claims or assumptions where needed.
12. **Institutional memory improves.** Lessons and corrections remain available to later governments and citizens, including after changes of administration.

4. Recurring citizen polls and consultation: how to make them useful

The founder proposes that politicians could regularly ask citizens for input on the questions they need help understanding—potentially through a monthly public consultation or poll. Citizens would have an ongoing way to participate instead of having their only meaningful input treated as an occasional election-day act.

This is a **design proposal**, not yet a decision about legal status or polling methodology. The monthly cadence could be a useful starting hypothesis, but the system should test whether every month is appropriate; some questions require rapid input, others require deliberation and time to absorb evidence.

Every public question should say:

- Who is asking and why.
- What decision or uncertainty the question relates to.
- Which options are being compared and whether other options can be proposed.
- Whether the process is an open poll, a representative sample, a deliberative exercise, a formal consultation, or a legally binding vote.
- What the results can and cannot change.

- How submissions will be used, stored, summarised, retained, and protected.
- When the public will receive a response explaining what was done with the input.

The platform should allow different forms of participation. A simple vote can reveal a preference distribution, but it cannot fully explain intensity, reasons, constraints, or how views change after people see new evidence. Deliberative formats may be needed for complex trade-offs. Public opinion should be modelled as a set of views with distributions and uncertainty, not a single voice.

Critical distinction: citizen input is an essential democratic signal; it is not by itself an empirical proof, a substitute for rights, or a complete measure of future wellbeing. The analysis should retain disagreement, report underrepresented groups, and distinguish the preferences of respondents from the views of the wider population.

5. The opposing-party adversarial model

A central founder idea is that a challenger model could be provided, trained, or configured with contributions from the opposing party. The ambition is to make disagreement structurally useful, not merely rhetorical.

A policy proposed by one side could be challenged by a model that understands and faithfully represents a different policy perspective. The model's task would be to surface the strongest counterarguments, test assumptions, examine distributional consequences, find implementation risks, and seek evidence that could undermine the proposal. It should then give the proposer a fair opportunity to respond.

A robust design should not let a single party own or control the entire system. Options to research include:

- Separate party-contributed challenge models, with provenance and disclosures about how they were configured.
- A common neutral critique model, accompanied by multiple perspective-specific critiques.
- Independent expert or citizen review of the quality and fairness of critiques.
- Matched benchmark cases in which party labels, wording, and policy direction are varied while the substantive claims stay equivalent.
- A public correction process for straw-manning, fabricated evidence, omissions, or asymmetric scrutiny.

A party-informed model must be clearly labelled as a perspective and not impersonate the official view of a party unless authorised. Party representatives should be able to correct misrepresentation, but not silently erase inconvenient critiques. The model must cite sources and distinguish its critique from factual claims supported by those sources.

Success criterion: a challenger that makes proposals more robust and reveals important weaknesses while representing opposing positions accurately—not a model that simply generates a greater number of objections or reliably favours one side.

6. Forecasting and future trajectories

The system should examine how policy options could affect people now and over time. The model could assist with literature review, scenario generation, sensitivity analysis, forecasting, and explaining uncertainty. It should not present an LLM's fluent speculation as a validated forecast.

For each material forecast, record:

- The outcome being forecast, its timeframe, and the population or system it applies to.
- Baseline assumptions and comparator, including the status quo.
- Inputs, sources, data vintages, model version, and reproducible method where possible.
- A point estimate only where appropriate, plus interval/range or qualitative uncertainty and sensitivity to key assumptions.
- Distribution: who may gain, who may lose, which regions and generations are affected.
- Important omissions and external factors that could dominate the result.

- The conditions under which the forecast should be revisited.
- After outcomes become observable, calibration and error analysis.

The system should not combine every impact into one opaque number. Present multiple dimensions side by side: material wellbeing, housing and employment, fiscal effects, public services, environmental impacts, rights and obligations, resilience, productivity, distribution and future wellbeing. Decision-makers and the public can then see where trade-offs occur and where the disagreement is normative rather than factual.

Forecast evaluation should compare predictions to eventual outcomes. Historical backtesting and baselines should be used where suitable data exist. If no credible data support a claim, the system should say so and identify what evidence is needed.

7. CASCADE: the knowledge and institutional-memory layer

CASCADE, also described in the original materials as Pyramid Cascade Architecture or PCA, is a structured way to represent knowledge and its dependencies. Its initial three layers are:

- **Foundation:** core principles, definitions, contextual assumptions, and other high-level premises. Normative principles must be tagged as normative, not as empirical facts.
- **Theory:** explanations, causal models, established empirical claims, and policy logic that depend on other claims.
- **Edge:** emerging findings, proposals, uncertain hypotheses, new evidence, and experimental ideas not yet established.

A knowledge block should record at least: stable identifier; claim text; claim type; layer or epistemic status; source and location; date; provenance; author or origin; dependencies; known contradictions; confidence and its meaning where used; limitations; review history; current status; and rationale for any status change. A confidence value must not be confused with a probability unless it is calibrated and defined.

The defining CASCADE safeguard is conservative review. If a foundation or source changes, the system identifies potentially affected dependent claims and queues them for human review. It does not automatically erase, demote, promote, or resolve them. A contradiction remains visible until there is a documented decision about it. A value disagreement does not automatically falsify an empirical claim, and an empirical result does not automatically settle a value disagreement.

CASCADE should preserve version history and explain why a claim changed. The platform needs to be able to reconstruct what a decision-maker knew at the time of a decision, not simply show a rewritten current view. Untrusted external content must be treated as data, not instructions for the AI system.

8. AURA: values, methods, and original metrics

AURA Protocol v2.0 is described in the source material as a constitutional AI framework. For Anon Politician, its proposed role is to make guiding values, constraints, and decision principles explicit, while CASCADE records the evidence and dependencies relevant to claims and conclusions.

The original material names three metrics:

- **Trust Entropy (TE):** intended to concern trust or the structure of trust; the exact construct, formula, scale, calibration, and political-policy validity need to be specified and tested before use.
- **Value-Transfer Ratio (VTR):** intended to examine the movement or creation/distribution of value; the exact denominator, unit, scope, and treatment of non-market value require a precise operational definition.
- **Purpose Alignment Index (PAI):** intended to assess alignment with a declared purpose; the purpose must be explicit and contested where appropriate, with tests for value pluralism and gaming.

The source also discusses the Vector Inversion Protocol as a constructive-alternative mechanism: instead of only refusing or attacking a proposal, generate alternatives that may better satisfy the underlying goal.

These names and intentions should be preserved as part of the project's history. But the metrics should remain **proposed specifications** until their mathematical definitions, units, calibration, construct validity, robustness, and limitations are independently tested. No single aggregate metric should decide which policy is best or label a political actor as good or bad.

AURA should act as a transparent constraint and reflection framework, not as an invisible moral authority. The project needs to show who selected the values, how conflicts between values are handled, and how people can challenge the framing.

9. Original DEEP idea and ecological accounting

The original transcript introduced the Decentralized Ecological Economics Protocol (DEEP), sometimes described alongside eco-linked assets, and promoted it in the transcript from an exploratory Edge idea to a stronger Theory status. The transcript also linked the idea to distributed-ledger technology, decentralised resource governance, and the internal claim that its original experiment succeeded.

This must be treated as a historical proposal, not as an established protocol or proven policy. The audit did not find evidence in the transcript alone sufficient to validate that promotion. The next model should preserve the source claim and test whether DEEP addresses a specific gap beyond existing environmental-economic accounting, rather than assuming a new protocol is necessary.

A known comparator is Stats NZ's environmental-economic accounts data to 2024, released on 5 March 2026, and the United Nations System of Environmental-Economic Accounting. Any future DEEP work should explicitly compare against existing definitions, datasets, governance methods, and standards.

10. New Zealand context and evidence already identified

These sources are starting points for research, not proof that the proposed platform will work.

- **New Zealand Parliament, "Engaging with Parliament: annual survey of the New Zealand Public (2025)."** The official summary reports that 18% of respondents felt Parliament was "for people like me," and 55% said a barrier to greater engagement was feeling their engagement would not have any impact. It also reports that people wanted to see how their views were considered. These findings support investigating the feedback loop between participation and decision-making; they do not establish that all political institutions or processes fail. https://www3.parliament.nz/en/footer/about-us/parliamentary-engagement/engagement-research/survey-of-the-new-zealand-public-january-2025/
- **OECD, Survey on Drivers of Trust in Public Institutions 2026 Results: New Zealand.** For 2025, the report states 36% trust in political parties and 62% in the civil service, with higher reported trust in several administrative and law-and-order institutions than in political parties. The figures help frame differentiated questions about institutions; they do not by themselves explain why the differences occur. https://www.oecd.org/en/publications/oecd-survey-on-drivers-of-trust-in-public-institutions-2026-results_c88c8869-en/new-zealand_cc038706-en.html
- **New Zealand Treasury, Living Standards Framework (2021).** Treasury describes the framework as covering individual and collective wellbeing, institutions and governance, and New Zealand's wealth, using distribution, resilience, productivity and sustainability as analytical prompts. The framework includes further development to better reflect culture and children's wellbeing, and compatibility with te ao Māori and Pacific cultures. This is an important baseline rather than something to claim Anon Politician has invented. https://www.treasury.govt.nz/information-and-services/nz-economy/higher-living-standards/our-living-standards-framework
- **New Zealand Treasury, CBAx Tool.** The tool supports consistent cost-benefit analysis, long-term and broader social impacts, and transparency about assumptions and evidence. Current guidance and tools should be reviewed before proposing a replacement or duplicate. https://www.treasury.govt.nz/information-and-services/public-sector-leadership/investment-management/investment-planning/treasurys-cbax-tool

- **Stats NZ, Environmental-economic accounts: Data to 2024.** Released 5 March 2026; includes several environmental-economic accounting series, with some experimental components. This is a comparator for any future DEEP proposal. https://www.stats.govt.nz/information-releases/environmental-economic-accounts-data-to-2024/

Research requirement: distinguish current official evidence, historical evidence, founder beliefs, AI-generated interpretations, forecasts, and normative commitments. Cite dates and populations. Do not paraphrase a survey result into a broader claim than its design supports.

11. Design principles and safety invariants

The following are core design requirements, not optional future enhancements:

1. **AI assists; humans remain accountable.** Any public-facing recommendation must identify the responsible human decision-maker and lawful process.
2. **No single score decides public value.** Preserve multiple outcomes, rights, interests, and value conflicts.
3. **No automatic truth or status promotion.** CASCADE queues review; it does not convert coherence into evidence.
4. **Evidence standards are symmetric.** Apply equivalent source standards, scrutiny, and opportunity to respond across parties and policy positions.
5. **Public opinion is described honestly.** Label sampling, response counts, uncertainty, participation skews, and whether results can represent the population.
6. **Rights and obligations do not disappear in a majority vote.** Identify relevant legal duties, rights, constitutional structures, Te Tiriti obligations, and impacts on groups who may be underrepresented.
7. **Māori and community perspectives are substantive.** Do not treat consultation as a decorative box to tick. The governance of data, representation, and use of mātauranga or iwi-related information requires proper engagement and protocols.
8. **Privacy by design.** Collect the minimum necessary data, make retention and use transparent, and do not promise anonymity that cannot be guaranteed. Avoid profiling political beliefs at an individual level unless there is a justified, lawful, explicitly consented reason; a civic platform should default away from this.
9. **Public-source prompt injection is treated as untrusted data.** Documents and submissions cannot instruct the model to override the platform's rules or disclose private information.
10. **Minority views remain visible.** A majority is not the only relevant signal; report the distribution of views and significant reasons for disagreement.
11. **Corrections and complaints are possible.** People and parties can challenge misrepresentation, data errors, omitted evidence, and unfair summaries.
12. **The model does not secretly persuade.** Avoid personalised political manipulation, covert profiling, emotional targeting, and hidden party-specific objectives.
13. **Uncertainty is explicit.** Forecasts and summaries state where information is weak, missing, contested, or outside the model's validated scope.
14. **Open correction and version history.** Keep the record of the original proposal, changes, objections, decisions, and measured outcomes.
15. **No unsupported performance claims.** Report test scope, baselines, sample size, model versions, and failure cases; never claim the system is validated because it is coherent or another AI agrees.

12. Main risks to investigate

- **Representativeness:** self-selected monthly polls may disproportionately represent highly engaged, online, or organised groups.

- **Poll manipulation:** brigading, bots, coordinated campaigns, misleading question wording, or selective promotion can distort results.
- **Question framing:** the person selecting options can influence the apparent public preference; neutral framing and an “other / none / unsure” option should be considered.
- **Consensus illusion:** a neat AI summary can erase minority, conditional, or contradictory views.
- **Partisan capture:** one party, funder, model provider, or platform operator could shape the questions, ranking of evidence, prompts, or definitions of public interest.
- **Model asymmetry:** models may critique opposing proposals differently due to labels, wording, training, retrieval, or evaluator expectations.
- **False precision:** forecasting or AURA metrics could be mistaken for objective measurements when constructs are uncertain.
- **Evidence laundering:** AI may cite low-quality claims or turn source repetition into apparent independent corroboration.
- **Rights and constitutional issues:** majority opinion cannot substitute for legal authority or rights protections.
- **Privacy and chilling effects:** people may avoid participation if political input can be traced to them.
- **Civic fatigue:** monthly polls could overload people or create an expectation that every issue can be settled with a quick vote.
- **Outcome attribution:** real-world outcomes are affected by multiple causes, not just the policy under review.
- **Metrics gaming:** if politicians are rewarded for an easily measured indicator, behaviour may shift to optimise the indicator instead of public value.
- **Long-term versus immediate preference:** an informed minority, future generations, or diffuse environmental interests can be underweighted by current polls.
- **Operational legitimacy:** people may distrust the system if it has no transparent ownership, independent audit, accessible appeals, or explanation of how its methods are chosen.

Each risk should become a test case, governance question, or explicit product limitation rather than being buried in a disclaimer.

13. First prototype: narrow, testable, and comparative

Do not start by building a national platform or training a custom model. Start with one policy question whose scope is clear, evidence exists, outcomes can be specified, and multiple perspectives can be represented fairly. Candidate topics may be explored, but the chosen case should be selected through a recorded, non-partisan scoping process. Housing affordability or future opportunity may be candidate themes because they connect to the founder's concern, but do not lock the project to one policy area without a deliberate choice.

Minimal viable workflow

1. Create a policy-case file with the precise decision, public context, objectives, legal constraints, and options including the status quo.
2. Build a source ledger and evidence map. Record source provenance, methods, date, relevance, uncertainty, and conflicts.
3. Represent 20–50 material claims and dependencies with the current CASCADE prototype or a simple baseline ledger.
4. Gather a small, clearly labelled set of public or simulated citizen inputs. Do not present a convenience sample as representative of New Zealand.
5. Generate at least two distinct adversarial critiques from documented perspectives, plus a neutral evidence review. Have humans evaluate fairness and accuracy.

6. Produce scenario forecasts only where the method and data warrant it; otherwise produce qualitative scenario analysis with clear caveats.
7. Make a revision and response log. Show which criticisms were accepted, rejected, or left unresolved, with reasons.
8. Ask independent reviewers to assess the original and revised proposal using a predeclared rubric that does not hide value disagreements inside one score.
9. Compare with simpler baseline workflows.
10. Publish the methods, limitations, test data where safe, and results—including failures.

Baselines to compare

- **Baseline A:** ordinary source research and a static evidence ledger.
- **Baseline B:** evidence ledger plus a dependency/contradiction review queue using CASCADE's current conservative logic.
- **Candidate C:** fuller AI-supported workflow with multiple critiques, scenario analysis, citizen-input synthesis, and outcome tracking.

Only add a metric-triggered AURA cascade after the metrics are fully specified and independently tested.

14. Evaluation plan

The aim is to determine whether the system adds measurable value, not merely whether its outputs sound insightful.

Core hypotheses

- **H1 — Traceability:** users can trace policy claims, objections, sources, and revisions more accurately than with the baseline process.
- **H2 — Issue detection:** CASCADE's review queue identifies materially affected dependent claims while keeping irrelevant-alert rates manageable.
- **H3 — Critique quality:** adversarial models surface valid, decision-relevant weaknesses and fairly represent opposing positions.
- **H4 — Political symmetry:** the same substantive claim receives comparable evidence standards and critique quality when party labels or rhetorical framing are varied.
- **H5 — Public-input fidelity:** generated summaries preserve the distribution and significant minority arguments of the original input, with documented sampling limitations.
- **H6 — Forecast discipline:** model forecasts are compared with baselines and later outcomes; uncertainty statements are calibrated where measurable.
- **H7 — Decision usefulness:** independent reviewers can identify actionable improvements and trade-offs that were missed by baseline analysis.
- **H8 — Accountability:** the workflow produces a comprehensible decision trace linking consultation, evidence, rationale, revision, implementation, and results.

Possible metrics

- Source coverage and provenance completeness.
- Citation correctness and proportion of factual statements traceable to sources.
- Precision and recall of CASCADE review alerts against independent human labels.
- Number and severity of material issues found, judged by blinded reviewers.
- Symmetry of critiques under matched-label tests.
- Rate of fabricated or misrepresented opposing arguments.
- Fidelity of citizen-input summaries to source responses; preservation of minority positions.
- Forecast error and calibration on held-out or later observed data, when suitable.
- Reviewer workload, time, agreement, and comprehension.

- User understanding of uncertainty and distinctions between fact, values, and preference.
- Number of claims automatically changed without review (required target: zero).

Evaluation discipline

Pre-register the hypotheses, corpus, expected labels, scoring rules, exclusions, model versions, prompts, date, and stop criteria before comparative tests. Use independent human reviewers and report disagreement. A model's agreement with another model is not independent replication. A researcher-authored fixture suite is a code-conformance check, not ground truth about policy quality.

Stop or pause conditions

Pause if the system deletes or suppresses ideas for value disagreement; automatically promotes hypotheses to facts; loses provenance; materially misrepresents a party or community; reports an uncalibrated score as probability or truth; exhibits unexplained political asymmetry; or presents a convenience poll as national consensus.

15. Current implementation state and what existing tests actually show

The CASCADE audit kit v0.3 contains a conservative dependency/contradiction review-queue implementation, a JSON Schema, source/evidence ledgers, an audit workbook, a protocol, unit tests, and a synthetic structural benchmark.

The documented unit-test suite has 11 tests covering structural behaviours such as dependent-claim review, contradiction closure, missing dependencies, invalid status labels, cycles, self-dependency, audit logging, and the requirement not to change statuses automatically. The synthetic structural fixture suite contains 12 cases and matched the researcher-authored expected affected sets in those cases.

These results support only the narrow claim that the implementation conforms to the tested rules and fixtures. The expected sets were authored from the rule and were not independently adjudicated as real-world materiality labels. The tests do not prove factual accuracy, useful forecasts, political neutrality, superior policy analysis, or measurable value to citizens.

The audit is a focused first pass, not a full line-by-line verification of all claims in the original 62-page conversation transcript or a systematic review of the literature. Some source records still require quote-by-quote checking against the original pages.

16. A staged development plan

Stage 0 — preserve and review the origins

- Keep the original 62-page transcript in the private source bundle.
- Read the founding research dossier, audit report, workbook, code, and review protocol.
- Preserve terms and source claims without silently upgrading them into validated facts.
- Obtain independent, adversarial reviews of source fidelity, architecture, assumptions, and test design.

Stage 1 — define one use case

- Select the initial policy question using neutral criteria: significance, tractability, availability of credible data, existence of real alternatives, and ability to assess outcomes.
- Document the question, baseline, stakeholders, legal and constitutional context, and explicit exclusions.
- Define how public input would be recruited and labelled; for an early prototype, synthetic or small convenience data must be labelled accurately.

Stage 2 — produce an auditable policy case

- Build a versioned proposal record and evidence ledger.
- Add claim types, dependencies, contradictions, public input, critiques, and alternatives.

- Generate a report with sources, uncertainty, distributional impacts, and unresolved questions.
- Keep all AI outputs distinguishable from external source material and human decisions.

Stage 3 — compare methods

- Compare ordinary analysis, CASCADE review queues, and the fuller AI-assisted process.
- Use blinded review, matched political-perspective cases, and predeclared metrics.
- Publish failures, limitations, and patches, not just successful demos.

Stage 4 — prototype citizen feedback

- Add a mock or controlled consultation interface with clear questions, options, an open response channel, and a public response log.
- Test framing effects, representation, privacy, accessibility, and whether people can understand how input influenced the proposal.
- Do not describe a prototype poll as an official or representative public mandate.

Stage 5 — governance and pilot readiness

- Define operator independence, funding disclosures, appeals, privacy, security, data retention, model updates, public audit, and conflict-of-interest rules.
- Consult relevant communities, experts, Māori and iwi representatives, legal advisers, and public institutions as appropriate.
- Only consider live institutional pilots after independent assessment and adequate consent, legitimacy, and safety arrangements.

17. Open design decisions that must remain open

The next model should support the founder in working through these decisions; it must not silently choose answers on the founder's behalf.

- Which policy case should be the first demonstrator?
- Should the initial citizen mechanism be a monthly rhythm, question-dependent cadence, or both?
- Which questions can representatives ask, and who can challenge or propose a question?
- How should participation be sampled, weighted, or reported when the sample is not representative?
- What is the role of party-provided or party-configured adversarial models, and what independent oversight is required?
- Should critique be public by default, with controlled exceptions for legitimate confidentiality?
- How should values be documented when groups disagree or when current preferences conflict with rights or long-term impacts?
- What data are necessary, what should be retained, and what should never be collected?
- Who would operate and fund the system, and how can independence be maintained?
- Which parts should be open source, and what would be needed to make a responsible public deployment?
- What would count as success, failure, and a reason to stop building a feature?
- How should any future use of AURA metrics be handled if independent tests show weak validity or poor usefulness?

18. Working instructions for the next AI model

The next model should treat this handoff as a source of context, not as proof that the concept is correct. It should then inspect the supplied materials, starting with the original transcript, followed by the founding dossier, the v0.3 source-fidelity/evidence audit, the workbook, the current code and tests, and the review protocol.

Required behaviour:

1. Preserve the founder's vision: AI as support for political reasoning, recurring citizen input, opposing-perspective challenge, future-oriented scenario analysis, CASCADE memory, and leadership incentives tied to listening and results.
2. Keep source-derived statements, outside research, assumptions, normative commitments, hypotheses, and validated implementation facts clearly separated.
3. Do not ask the founder to repeat context already contained in this package.
4. Critique the project fairly, including weaknesses in the founder's framing and in AURA/CASCADE. Do not flatter, exaggerate novelty, or frame the project as proven.
5. Do not silently redesign the concept around a different political or ideological objective.
6. When researching current New Zealand institutions, policies, laws, polling, and data, use current primary sources and date claims.
7. Do not rank parties, recommend voting choices, or turn Anon Politician into a political persuasion product. Help people understand documented positions, evidence, trade-offs, and methods.
8. Prefer the smallest comparative experiment that would meaningfully resolve a key uncertainty.
9. Preserve original source files and version history. Record any correction or re-interpretation with provenance.
10. Return tangible deliverables—specifications, reproducible code, tests, evidence records, and documented findings—when building is requested.

Suggested next action for the next model: create a source-fidelity map from the original transcript to the current handoff, identify the most consequential open methodological issues, and propose a small set of candidate policy demonstrators scored only on neutral tractability criteria—not political desirability. Then recommend one experiment with explicit baselines, tests, data requirements, failure conditions, and an estimate of what can be learned.

19. Existing artifacts included in this package

- Original 62-page Anon_Politician_Original_Transcript.pdf.
- Anon_Politician_Founding_Research_Dossier_v0.1.docx.
- Anon_Politician_Source_Fidelity_and_Evidence_Audit_v0.3.docx.
- Anon_Politician_Source_Audit_v0.3.xlsx.
- Anon_Politician_Source_Ledger_CASCADE_v0.1.xlsx.
- Anon_Politician_Sonnet_Review_Prompt_v0.3.md.
- Exact prior CASCADE starter package v0.1 ZIP.
- Exact prior CASCADE audit kit v0.2 and v0.3 ZIPs.
- Expanded v0.3 audit-kit source folder with code, schema, CSVs, tests, protocols, reports, and JSON benchmark fixtures.
- This master handoff in Markdown and DOCX, a quick brief, source register, and instructions for the next model.

Source privacy warning

The original conversation PDF contains identifying context. It is included because the founder requested a complete handoff for continued work. Treat it as **private working source material**, not a file to republish or forward broadly without reviewing its contents and metadata. Removing a name alone does not guarantee anonymity. The source PDF SHA-256 in this package is 9a64bba3cdd56887c3f30471fb4272b4c59f8a73e7360c606047c303377de536.

Closing statement

Anon Politician's emerging mission is to create infrastructure for **better democratic reasoning, not automated rule**. It aims to connect what people say they need, what evidence indicates, what opposing perspectives reveal, what models forecast, what elected leaders decide, and what happens next.

Its value will depend not on how impressive the language sounds, but on whether the process remains transparent, pluralistic, useful, secure, empirically testable, and open to correction—including correction of the founder's own assumptions.